



PostgreSQL as an integrated data analysis platform

FOSS4G.be 2015 – Oslandia Team

Let 's try to Think Different (about PostgreSQL)

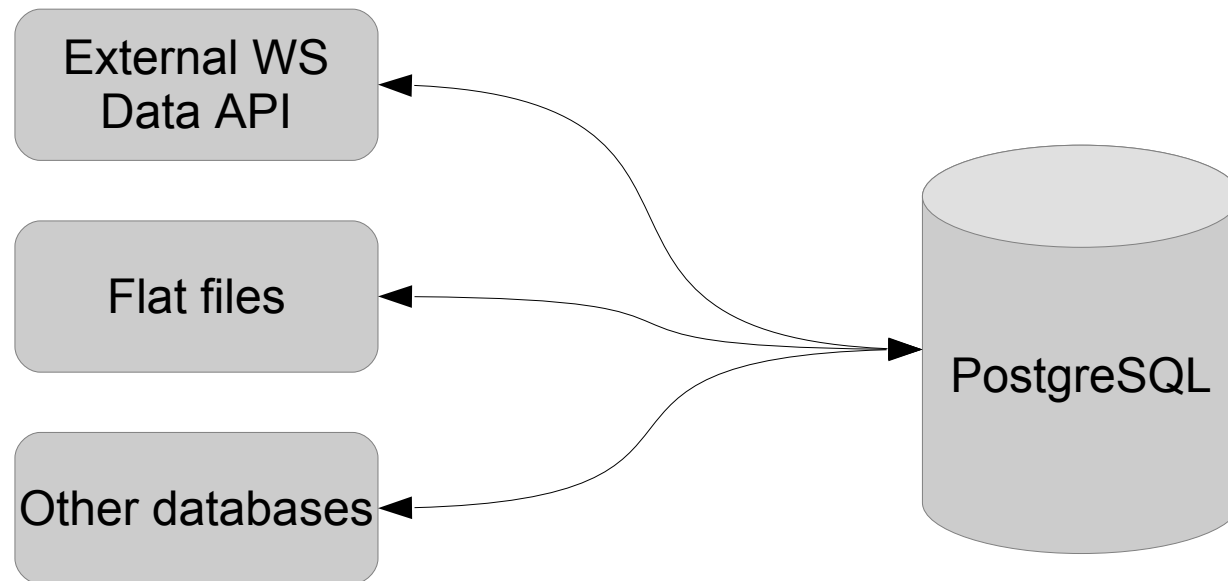
Database is not only a place to store data
(and use basic SQL to access it)

Let 's try to Think Different (about PostgreSQL)

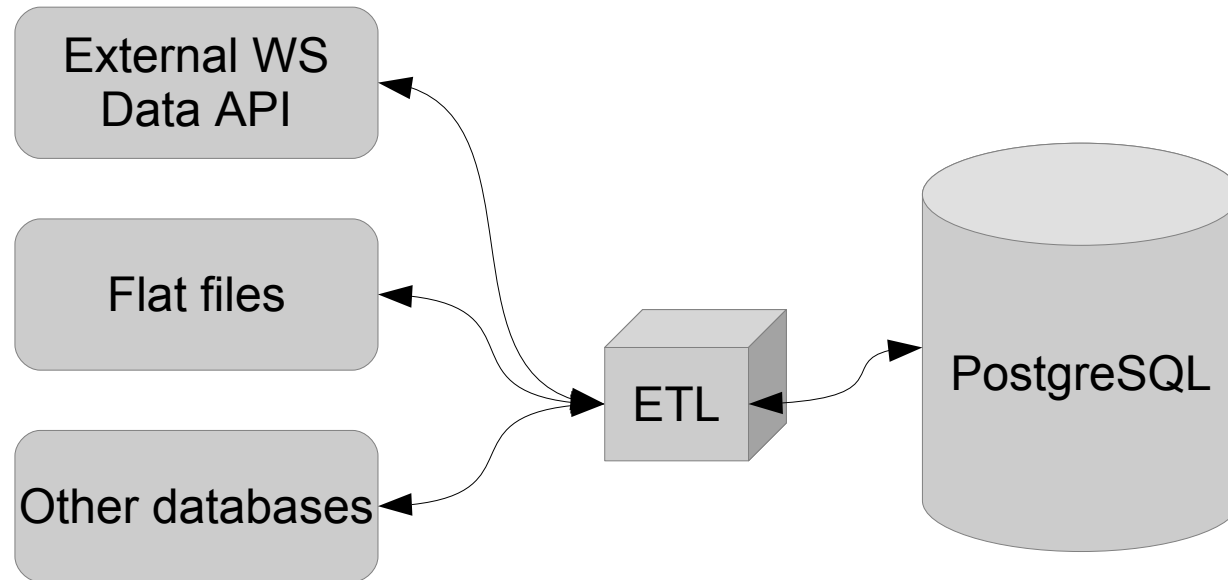
Database is not only a place to store data
(and use basic SQL to access it)

PostgreSQL is far more than an enhanced filesystem
PostgreSQL by design is extensible

#1 Data Integration

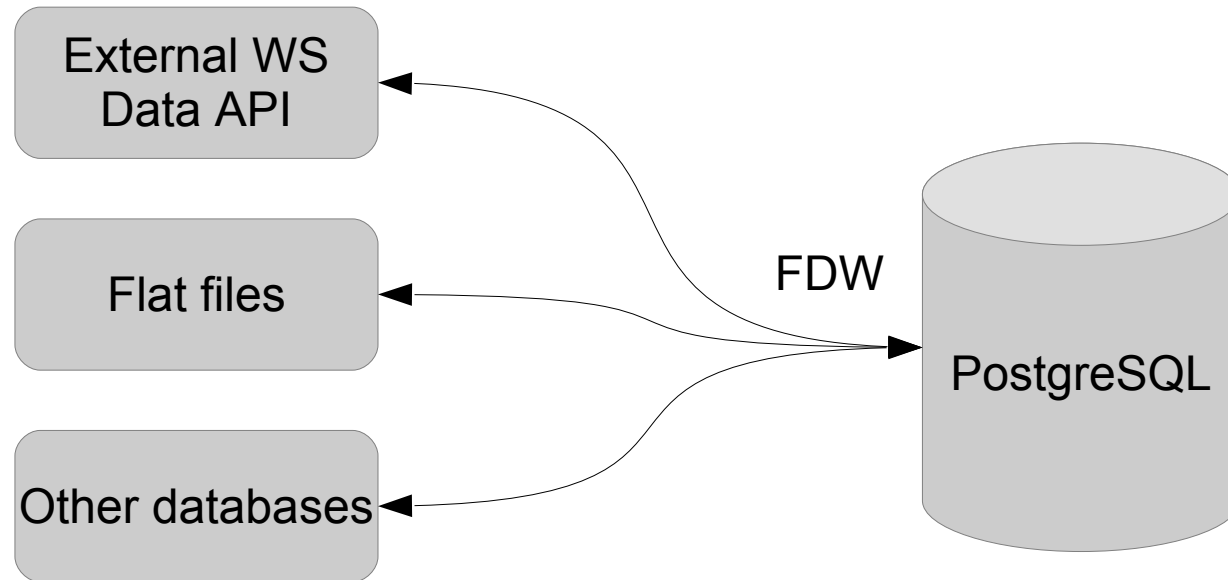


#1 Data Integration



Common answer is « Use an ETL »

#1 Data Integration




Alternate answer is
« Use PostgreSQL Foreign Data Wrapper »

https://wiki.postgresql.org/wiki/Foreign_data_wrappers

SQL Management of External Data (SQL/MED) added to the SQL standard
Handling access to remote objects from SQL databases
Available in PostgreSQL since 9.3

#1 Data integration : OGR FDW



https://wiki.postgresql.org/wiki/Foreign_data_wrappers

SQL Management of External Data (SQL/MED) added to the SQL standard
Handling access to remote objects from SQL databases
Available in PostgreSQL since 9.3

- 1 Foreign Data Wrappers
 - 1.1 Generic SQL Database Wrappers
 - 1.2 Specific SQL Database Wrappers
 - 1.3 NoSQL Database Wrappers
 - 1.4 File Wrappers
 - 1.5 Geo Wrappers
 - 1.6 LDAP Wrappers
 - 1.7 Generic Web Wrappers
 - 1.8 Specific Web Wrappers
 - 1.9 Big Data Wrappers
 - 1.10 Column-Oriented Wrappers
 - 1.11 Scientific Wrappers
 - 1.12 Operating System Wrappers
 - 1.13 Exotic Wrappers
 - 1.14 Example Wrappers

~50 native connectors already available
(And more through Multicorn extension)

<https://github.com/pramsey/pgsql-ogr-fdw>

Install OGR FDW

```
git clone https://github.com/pramsey/pgsql-ogr-fdw.git
cd postgresql-ogr-fdw
make
sudo make install
```

Define a FDW wrapper

```
CREATE EXTENSION postgis;  
CREATE EXTENSION ogr_fdw;  
  
CREATE SERVER shapefile_france  
  FOREIGN DATA WRAPPER ogr_fdw  
  OPTIONS (  
    datasource '/tmp/fdw_ogr/france.shp',  
    format 'ESRI Shapefile'  
  );
```

Retrieve shapefile attributes list (metadata)

```
ogrinfo -al -so /tmp/fdw_ogr/france.shp
```

Create Foreign table

```
CREATE SCHEMA shp;  
  
CREATE FOREIGN TABLE shp.france (  
    id_geofla integer,  
    geom geometry,  
    code_chf_l varchar,  
    nom_chf_l varchar,  
    x_chf_lieu varchar,  
    y_chf_lieu varchar,  
    x_centroid integer,  
    y_centroid integer,  
    nom_dept varchar,  
    code_reg varchar,  
    nom_region varchar,  
    code_dept varchar  
)  
  
SERVER shapefile_france  
OPTIONS (layer 'france');
```

Check it

```
SELECT id_geofla, ST_AsEWKT(ST_Centroid(geom)) AS geom  
FROM shp.france LIMIT 1 ;
```

Create VIEW from Foreign Table

<https://github.com/pramsey/pgsql-ogr-fdw/issues/11>

```
CREATE OR REPLACE VIEW shp.france_wfs AS

SELECT id_geofla,
       ST_Multi(ST_SetSRID(geom,27572))::geometry(MultiPolygon,27572) AS geom,
       code_dept,
       nom_dept
FROM france;
```

TinyOWS configuration

```
<tinyows online_resource="http://127.0.0.1/cgi-bin/tinyows"
          schema_dir="/usr/local/share/tinyows/schema/"
          estimated_extent="1"
          display_bbox="0">

  <pg host="127.0.0.1" user="pggis" password="***" dbname="db" />

  <metadata name="TinyOWS WFS Server"
            title="TinyOWS Server – OGR FDW Service" />

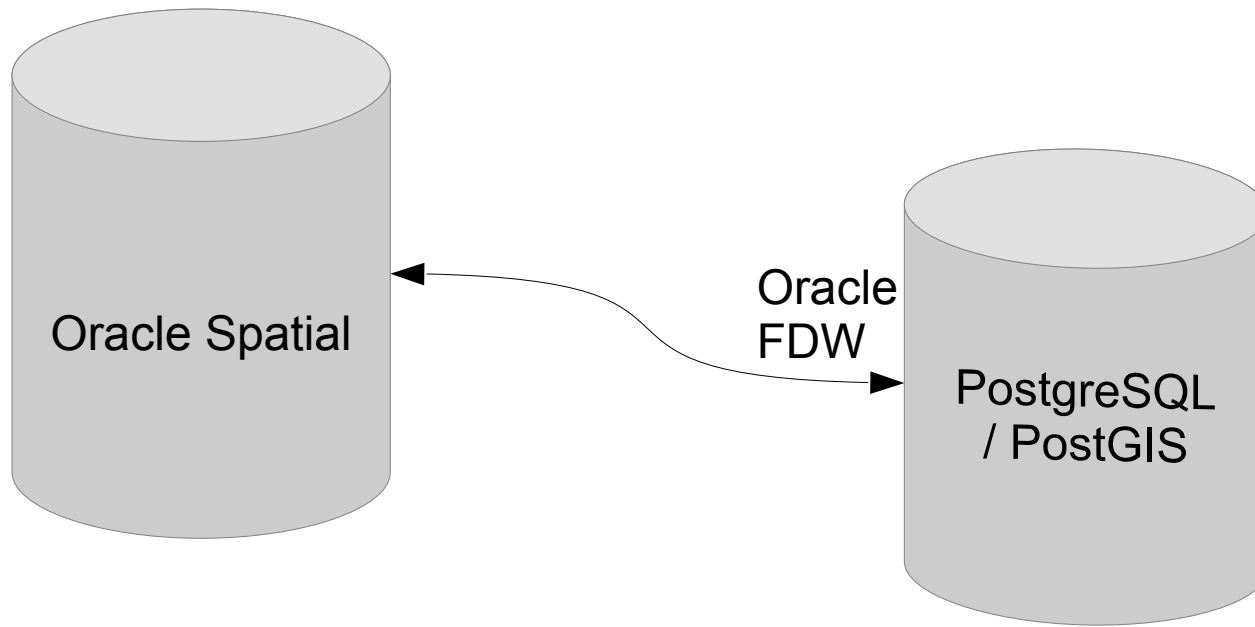
  <layer retrievable="1"
        writable="0"
        ns_prefix="tows"
        ns_uri="http://www.tinyows.org/"
        schema="shp"
        name="france_wfs"
        title="france" />

</tinyows>
```

Check it

```
wget -O out http://127.0.0.1/cgi-bin/tinyows?
SERVICE=WFS&REQUEST=GetFeature&Typename=tows:france_wfs
```

#1 Data integration : Oracle FDW



http://pgxn.org/dist/oracle_fdw/

```
CREATE EXTENSION postgres_fdw;  
CREATE EXTENSION oracle_fdw;  
  
CREATE SERVER orcl FOREIGN DATA WRAPPER oracle_fdw  
OPTIONS (dbserver '${ORACLE_URI}');
```

Oracle user Mapping

```
GRANT USAGE ON FOREIGN SERVER orcl TO ${PGUSER};  
  
CREATE USER orcl_map FOR ${PGUSER}  
SERVER orcl  
OPTIONS (user '${ORAUSER}', password '${ORAPWD}');
```



```
CREATE SCHEMA fdw;
```

```
CREATE FOREIGN TABLE fdw.foo (  
    id double precision,  
    label varchar,  
    last_update date,  
    geom geometry(POINT, 2154),  
)  
SERVER orcl  
OPTIONS (schema '${ORAUSER}', table 'F00');
```

```
CREATE SCHEMA mat;  
CREATE MATERIALIZED VIEW mat.foo AS SELECT * FROM fdw.foo;
```

```
CREATE UNIQUE INDEX ON mat.foo(id);  
CREATE INDEX ON mat.foo USING GIST(geom);
```


```
REFRESH MATERIALIZED VIEW CONCURRENTLY mat.foo;
```

#2 Cleaning Data: Validity


```
SELECT count(*) FROM my_schema.my_table WHERE NOT ST_IsValid(geom);
```

#2 Cleaning Data: Validity


```
SELECT count(*) FROM my_schema.my_table WHERE NOT ST_IsValid(geom);
```




(1)
Hole crosses shell




(2)
Hole touches shell at
more than one point




(3)
Hole touches shell in
line segment




(4)
The polygon interior is
disconnected



(5)
Holes cross



(6)
Holes touch in line
segment



(7)
Shell self-intersects

```
UPDATE my_schema."my_table"
```

```
SET geom=ST_CollectionExtract(ST_MakeValid(geom), 3)
```

```
WHERE ST_IsValidReason(geom) != 'Valid Geometry'  
      AND (GeometryType(geom) = 'POLYGON'  
           OR GeometryType(geom) = 'MULTIPOLYGON');
```

```
UPDATE my_schema."my_table"
```

```
SET geom=ST_CollectionExtract(ST_MakeValid(geom), 3)
```

```
WHERE ST_IsValidReason(geom) != 'Valid Geometry'  
      AND (GeometryType(geom) = 'POLYGON'  
          OR GeometryType(geom) = 'MULTIPOLYGON');
```

Still to deal with :


- Null Surface → Empty
- Single Point Line → infinitesimal ending point translation

#2 Cleaning Data: Reclassify

```
SELECT  id, input,  
        CASE WHEN input = 'yes' THEN 1::boolean  
              WHEN input = 'no'  THEN 0::boolean  
              ELSE NULL  
        END reclass  
FROM data;
```

id	input	reclass
1	yes	t
2	no	f
3	NC	

#3 Data Processing



Common answer is :
« Develop an external script »

#3 Data Processing



Alternate answer is :
«Hey it's already there !»

Since PostgreSQL 9.1 : EXTENSION handling

Using existing extension is that easy, UUID generation example :

```
foo=# CREATE EXTENSION "uuid-oss";  
CREATE EXTENSION
```

```
foo=# SELECT uuid_generate_v4();
```

```
6953879c-3aae-4d42-a470-6d430305e173
```

Lot of PostgreSQL extensions available (really)

To display those already available on your server :

```
SELECT * FROM pg_available_extensions ;
```

An PostgreSQL extension repository:

<http://pgxn.org/>

PGXN, the PostgreSQL Extension network, is a central distribution system for open-source PostgreSQL extension libraries.

Recent Releases

pg_shard 1.2.2
Easy sharding for PostgreSQL

test_factory 0.3.1
Framework for managing test data

test_factory 0.3.0
Framework for managing test data

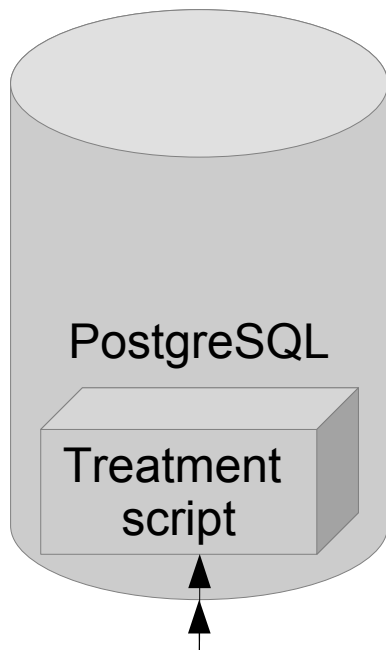
hyperloglog_estimator 1.2.6
Estimates number of distinct elements in a data set (aggregate and a data type).

geoip 0.2.4
Geolocation using GeoIP.

Some useful PostgreSQL extensions (among others)

- `pg_trgm`
Use trigram matching to evaluate string similarity (for natural language texts search)
- `Fuzzystrmatch`
Alternates well known string similarity functions (levenshtein, soundex...)
- `Unaccent`
Deal with accentuated text
- `xml2`
Xpath functions facilities (use `libxml2`)
- `Pgcrypto`
Cryptographic functions
- `Hstore`
Storing and manipulation of key/value pairs inside a single PostgreSQL value

#3 Data Processing



Alternate answer is :
«Put your scripts inside PostgreSQL»

#3 Data Processing : PL/Python

Using existing Python Library from PostgreSQL
Through SQL function

#3 Data Processing : PL/Python

Using existing Python Library from PostgreSQL
Call through SQL function

An example with GeoPy, Installation :

```
sudo apt-get install postgresql-plpython-9.4 python3-geopy  
  
createdb db  
createlang plpython3u db  
psql db -c "CREATE EXTENSION postgis"
```

Register on GeoNames

Enable your account to use the free WebService

PI/Python basic Geocoder function

```
CREATE OR REPLACE FUNCTION geoname(toponym text)
    RETURNS geometry(Point,4326)
AS $$

    from geopy import geocoders
    g = geocoders.GeoNames(username="YOUR_USERNAME")

    try:
        place, (lat, lng) = g.geocode(toponym)

        result = plpy.execute(
            "SELECT 'SRID=4326;POINT(%s %s)::geometry(Point, 4326) AS geom"
            % (lng, lat), 1)

        return result[0]["geom"]


    except:
        plpy.warning('Geocoding Error')
        return None

$$ LANGUAGE plpython3u;
```

Check it :

```
psql db -c  
"SELECT ST_AsGeoJSON(geoname('New York, NY 10022'))"  
  
{"type":"Point","coordinates":[-74.00597,40.71427]}
```

<http://www.openstreetmap.org/?mlon=-74.00597&mlat=40.71427&zoom=12>



#3 Data Processing : GeoSpatial statistic correlation

```
WITH lyon AS (SELECT ST_Transform(geoname('Lyon'), 2154) geom),  
    tc AS (SELECT DISTINCT ON (tc.gid) tc.gid, v.gid, tc.geom,  
        ST_Distance(tc.geom, v.geom) dist  
        FROM data.tc tc, data.velov v  
        ORDER BY tc.gid, tc.geom <-> v.geom)  
  
SELECT corr(dist, ST_Distance(l.geom, tc.geom)) FROM lyon l, tc;
```


#3 Data Processing : GeoSpatial statistic correlation

```
WITH lyon AS (SELECT ST_Transform(geoname('Lyon'), 2154) geom),  
    tc AS (SELECT DISTINCT ON (tc.gid) tc.gid, v.gid, tc.geom,  
        ST_Distance(tc.geom, v.geom) dist  
        FROM data.tc tc, data.velov v  
        ORDER BY tc.gid, tc.geom <-> v.geom)  
  
SELECT corr(dist, ST_Distance(l.geom, tc.geom)) FROM lyon l, tc;
```

corr

0.951847587972223

#3 Data Processing : GeoSpatial statistic correlation



#3 Data Processing : Same, but via PL/R

```
createlang plr DATABASE
```

```
CREATE OR REPLACE FUNCTION r_corr(a float[], b float[])  
  RETURNS float AS  
$$  
  return (cor(a, b))  
$$ language plr;
```

```

WITH lyon AS (SELECT ST_Transform(geoname('Lyon'), 2154) geom),
  tc AS (SELECT DISTINCT ON (tc.gid) tc.gid,
    ST_Distance(tc.geom, v.geom) dist, tc.geom
    FROM data.tc tc, data.velov v ORDER BY tc.gid, tc.geom <-> v.geom)

SELECT r_corr(ARRAY(SELECT dist
    FROM lyon l, tc ORDER BY tc.gid),
    ARRAY(SELECT(ST_Distance(l.geom, tc.geom))
    FROM lyon l, tc ORDER BY tc.gid)
);

```

r_corr

0.951847587972213

#To Go Further

Write his own PostgreSQL module in C (if needed)

Write your own FDW (with Multicorn or in C)

Use PL/R to use R advanced statistics

Machine Learning

#Conclusion

PostgreSQL behaves like an extensible and integrated Framework

Allow you to keep all data in the same place

(modern) SQL acting as a glue language



www.oslandia.com

Thanks !